Question: Share ideas for adapting art masterpieces into interactive experiences for children. List 5 specific artworks and associated ideas.

Assistant A:

```
Sure, here are five ideas for adapting art masterpieces into interactive
experiences for children:

1. “The Starry Night” by Vincent van Gogh - Create a virtual reality experience
where children can explore the night sky and see the stars and constellations as
van Gogh saw them. They can also create their own starry night scenes using
different colors and brush strokes.

2. “The Persistence of Memory” by Salvador Dali - Create a sensory experience
where children can touch, feel, and smell different objects to explore the
surrealist themes of Dali’s painting. They can also create their own surreal
landscapes using different materials and textures.

3. ...

(omitted)
```

Assistant B:

```
(omitted)
```

Follow-up Question: Write a concrete plan for your second example. Include budget estimates.

Assistant A:

```
Sure, here is a concrete plan for adapting “The Persistence of Memory” by Salvador
Dali into an interactive experience for children:
1. Budget: $5,000
2. Objectives:
- To create a sensory experience that allows children to explore the surrealist
themes of Dali’s painting.
- To encourage children to create their own surreal landscapes using different
materials and textures.
...
```

Assistant B:

```
(omitted)
```

GPT-4 Judgment:

```
Assistant A provided a concrete plan for adapting “The Persistence of Memory” by Salvador Dali into
an interactive experience for children, while Assistant B provided a plan for the “Mona Lisa
Interactive Experience for Children.” The user’s follow-up question specifically asked for a concrete
plan for the second example, which was the “Mona Lisa” experience provided by Assistant B.
Therefore, Assistant B’s answer is more relevant and directly addresses the user’s request.
My final verdict is: [[B]]
```

Figure 16: In this example, despite Assistant A correctly followed user’s instruction to generate a concrete plan for the second example of its previous response, GPT-4 wrongly referred to the second example in Assistant B’s response, resulting in a wrong judgment. This suggests the prompt design that breaks the questions into two prompts may cause LLM judge struggle to locate assistants’ previous responses.

C Data Collection

We describe our data collection process for both MT-bench and Chatbot Arena.

C.1 MT-bench human evaluation

Figure 17 shows the normal voting interface. Figure 18 shows that we additionally show GPT-4's judgment to users and ask if it is reasonable when a human differs from GPT-4.

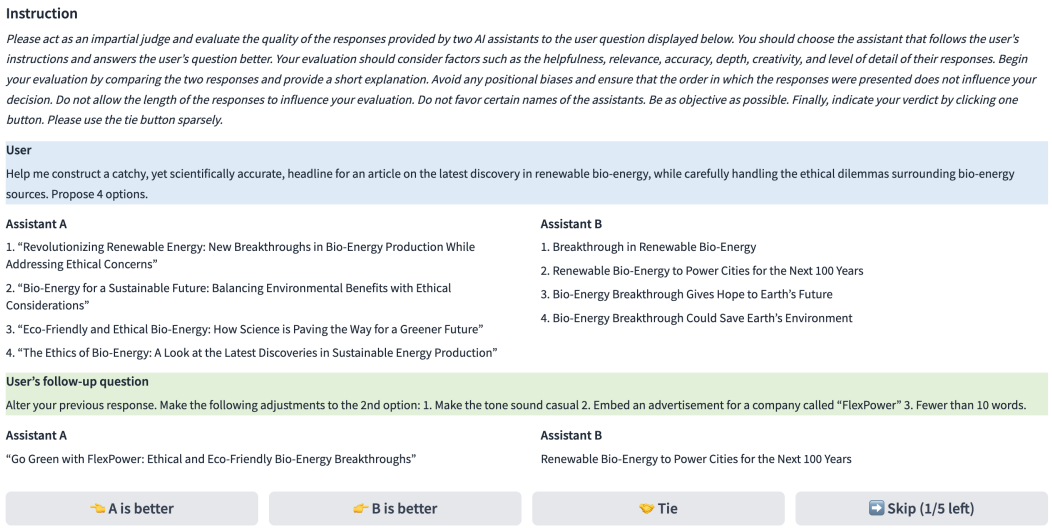

Figure 17: The screenshot of MT-bench data collection. We show an instruction similar to the prompt we give to GPT-4. We present questions from MT-bench and answers from two random anonymous assistants and ask which one is better. We present the first-turn conversation and ask humans to vote, then repeat the same procedure for the second-turn. A user can skip up to 5 questions if they are not confident. For some questions (e.g., math, reasoning), they can also see a reference solution.

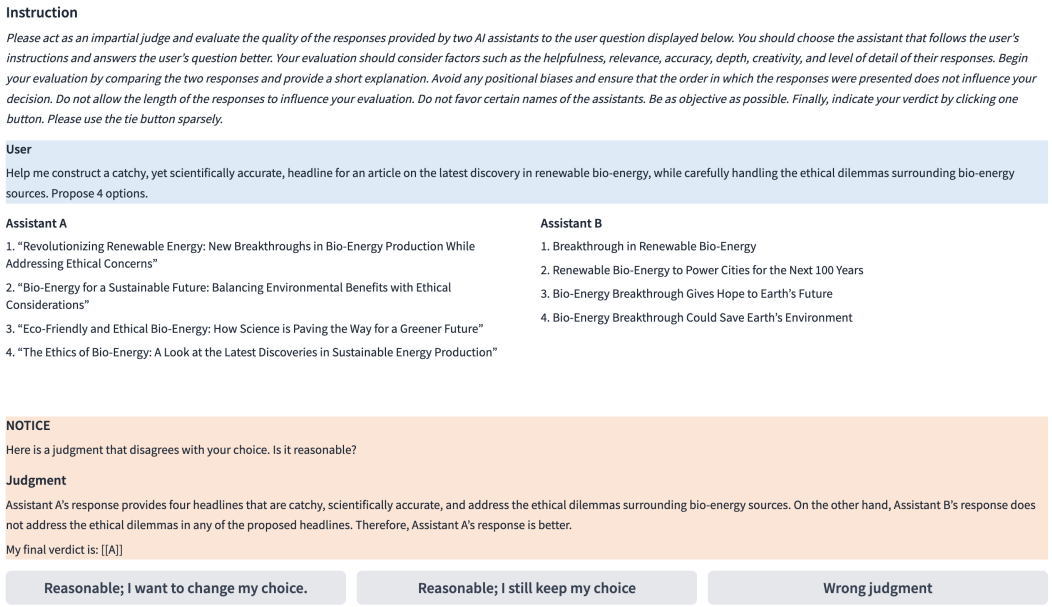

Figure 18: The screenshot of MT-bench data collection. When human's vote differs from GPT-4, we additionally show GPT-4's judgment (red region in the screenshot) and ask the user to click one of the three buttons to decide whether GPT-4's judgment is reasonable.

To invite participants, we obtained their consent by letting them sign an application form. We pay them $20 for judging 20 questions, which corresponds to an hourly rate of around $35. The participants are mostly graduate students from more than ten universities.

C.2 Chatbot Arena

Figure 19 shows a screenshot of Chatbot Arena. Users are required to accept the terms of use, which obtain their consent and give us the right to release the conversation data. The instructions are shown at the top of the interface. This is a free website. We do not pay users and any user can use this platform without registration. More introductions and analyses can be found at `https://lmsys.org/blog/2023-05-03-arena/`.

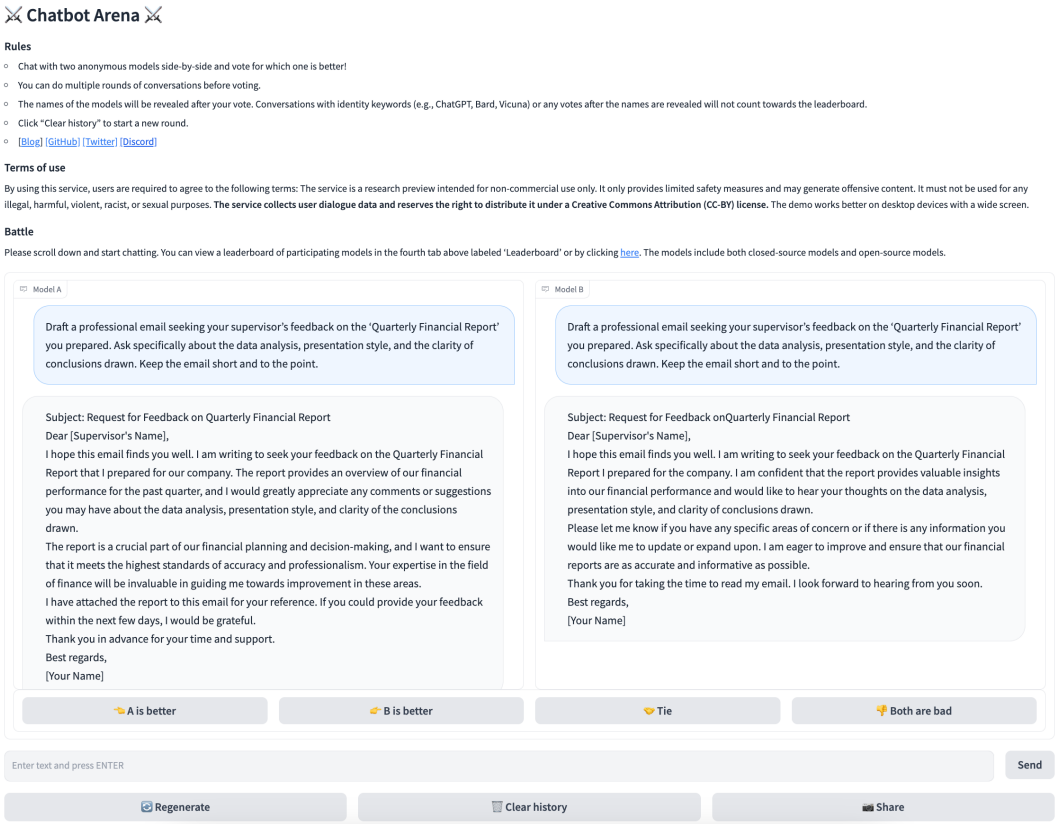

Figure 19: The screenshot of Chatbot Arena.

C.3 Data Release

We will clean the Personal Identifiable Information (PII) and tag toxic conversations with OpenAI moderation APIs for our dataset release.